

# MidoNet Quick Start Guide

for RHEL 7 / Kilo (OSP)

5.0-SNAPSHOT (2015-11-12 07:28 UTC)

DRAFT



## MidoNet Quick Start Guide for RHEL 7 / Kilo (OSP)

5.0-SNAPSHOT (2015-11-12 07:28 UTC)

Copyright © 2015 Midokura SARL All rights reserved.

MidoNet is a network virtualization software for Infrastructure-as-a-Service (IaaS) clouds.

It decouples your IaaS cloud from your network hardware, creating an intelligent software abstraction layer between your end hosts and your physical network.

This guide walks through the minimum installation and configuration steps necessary to use MidoNet with OpenStack.



### Caution

This document is a DRAFT. It may be MISSING relevant information or contain UNTESTED information. Use it at your own risk.



### Note

Please consult the [MidoNet Mailing Lists or Chat](#) if you need assistance.

Licensed under the Apache License, Version 2.0 (the "License"); you may not use this file except in compliance with the License. You may obtain a copy of the License at

<http://www.apache.org/licenses/LICENSE-2.0>

Unless required by applicable law or agreed to in writing, software distributed under the License is distributed on an "AS IS" BASIS, WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or implied. See the License for the specific language governing permissions and limitations under the License.

## Table of Contents

Preface .....	iv
Conventions .....	iv
1. Architecture .....	1
Hosts and Services .....	2
2. Basic Environment Configuration .....	4
Networking Configuration .....	4
SELinux Configuration .....	4
Repository Configuration .....	4
3. OpenStack Installation .....	6
Identity Service (Keystone) .....	6
Compute Services (Nova) .....	6
Networking Services (Neutron) .....	10
4. MidoNet Installation .....	14
NSDB Nodes .....	14
Controller Node .....	17
Midolman Installation .....	19
MidoNet Host Registration .....	20
5. BGP Uplink Configuration .....	22
6. Further Steps .....	26

# Preface

## Conventions

The MidoNet documentation uses several typesetting conventions.

## Notices

Notices take these forms:



### Note

A handy tip or reminder.



### Important

Something you must be aware of before proceeding.



### Warning

Critical information about the risk of data loss or security issues.

## Command prompts

### \$ prompt

Any user, including the root user, can run commands that are prefixed with the \$ prompt.

### # prompt

The root user must run commands that are prefixed with the # prompt. You can also prefix these commands with the **sudo** command, if available, to run them.

# 1. Architecture

## Table of Contents

Hosts and Services .....	2
--------------------------	---

This guide assumes the following example system architecture.

OpenStack Controller Node:

- Controller Node (**controller**)

Compute Node:

- Compute Node (**compute1**)

Since MidoNet is a distributed system, it does not have the concept of a Network Node as being used with the default OpenStack networking plugin. Instead it uses two or more Gateway Nodes that utilize [Quagga](#) to provide connectivity to external networks via the Border Gateway Protocol (BGP).

- Gateway Node 1 (**gateway1**)
- Gateway Node 2 (**gateway2**)

Three or more hosts are being used for the MidoNet Network State Database (NSDB) cluster which utilizes [ZooKeeper](#) and [Cassandra](#) to store virtual network topology and connection state information:

- NSDB Node 1 (**nsdb1**)
- NSDB Node 2 (**nsdb2**)
- NSDB Node 3 (**nsdb3**)



### Important

Ideally, both the ZooKeeper transaction log and Cassandra data files need their own dedicated disks, with additional disks for other services on the host. However, for small POCs and small deployments, it is ok to share the Cassandra disk with other services and just leave the ZooKeeper transaction log on its own.

The *MidoNet Agent (Midolman)* has to be installed on all nodes where traffic enters or leaves the virtual topology. In this guide this are the **controller**, **gateway1**, **gateway2** and **compute1** hosts.

The *Midonet Cluster* can be installed on a separate host, but this guide assumes it to be installed on the **controller** host.

The *Midonet Command Line Interface (CLI)* can be installed on any host that has connectivity to the MidoNet Cluster. This guide assumes it to be installed on the **controller** host.

The *Midonet Neutron Plugin* replaces the ML2 Plugin and has to be installed on the **controller**.

## Hosts and Services

### Controller Node (controller)

- General
  - Database (MariaDB)
  - Message Broker (RabbitMQ)
- OpenStack
  - Identity Service (Keystone)
  - Image Service (Glance)
  - Compute (Nova)
  - Networking (Neutron)
    - Neutron Server
    - DHCP Agent
    - Metadata Agent
  - Dashboard (Horizon)
- MidoNet
  - Cluster
  - CLI
  - Neutron Plugin

### Compute Node (compute1)

- OpenStack
  - Compute (Nova)
  - Networking (Neutron)
- MidoNet
  - Agent (Midolman)

### NSDB Nodes (nsdb1, nsdb2, nsdb3)

- Network State Database (NSDB)
  - Network Topology (ZooKeeper)
  - Network State Information (Cassandra)

### Gateway Nodes (gateway1, gateway2)

- BGP Daemon (Quagga)
- MidoNet
  - Agent (Midolman)

## 2. Basic Environment Configuration

### Table of Contents

Networking Configuration .....	4
SELinux Configuration .....	4
Repository Configuration .....	4

### Networking Configuration



#### Important

All hostnames must be resolvable, either via DNS or locally.

This guide assumes that you follow the instructions in [OpenStack Networking \(neutron\)](#) of the OpenStack Documentation.

### SELinux Configuration



#### Important

This guide assumes that SELinux (if installed) is either in permissive state or disabled.

To change the mode, execute the following command:

```
# setenforce Permissive
```

To permanently change the SELinux configuration, edit the `/etc/selinux/config` file accordingly:

```
SELINUX=permissive
```

### Repository Configuration

Configure necessary software repositories and update installed packages.

#### 1. Enable Red Hat base repository

```
# subscription-manager repos --enable=rhel-7-server-rpms
```

#### 2. Enable Red Hat OSP repository

```
# subscription-manager repos --enable=rhel-7-server-openstack-7.0-rpms
```

#### 3. Enable DataStax repository

Create the `/etc/yum.repos.d/datastax.repo` file and edit it to contain the following:

```
# DataStax (Apache Cassandra)
[datastax]
name = DataStax Repo for Apache Cassandra
```



```
baseurl = http://rpm.datastax.com/community
enabled = 1
gpgcheck = 1
gpgkey = https://rpm.datastax.com/rpm/repo_key
```

### 1. Enable MidoNet repositories

Create the `/etc/yum.repos.d/midonet.repo` file and edit it to contain the following:

```
[midonet]
name=MidoNet
baseurl=http://builds.midonet.org/midonet-5/stable/el7/
enabled=1
gpgcheck=1
gpgkey=http://builds.midonet.org/midorepo.key

[midonet-openstack-integration]
name=MidoNet OpenStack Integration
baseurl=http://builds.midonet.org/openstack-kilo/stable/el7/
enabled=1
gpgcheck=1
gpgkey=http://builds.midonet.org/midorepo.key

[midonet-misc]
name=MidoNet 3rd Party Tools and Libraries
baseurl=http://builds.midonet.org/misc/stable/el7/
enabled=1
gpgcheck=1
gpgkey=http://builds.midonet.org/midorepo.key
```

### 2. Install available updates

```
# yum clean all
# yum upgrade
```

### 3. If necessary, reboot the system

```
# reboot
```

### 3. OpenStack Installation

# Table of Contents

Identity Service (Keystone) .....	6
Compute Services (Nova) .....	6
Networking Services (Neutron) .....	10



## Important

Follow the [Installation Reference](#) documentation, but **note the following differences**.

## Identity Service (Keystone)



## Important

Follow the Red Hat documentation's [Chapter 3. Install The Identity Service](#) instructions, but **note the following additions**.

## 1. Create MidoNet API Service

As Keystone admin, execute the following command:

```
$ openstack service create --name midonet --description "MidoNet API Service" midonet
```

## 2. Create MidoNet Administrative User

As Keystone admin, execute the following commands:

```
$ keystone user-create --name midonet --pass MIDONET_PASS --tenant services
$ keystone user-role-add --user midonet --role admin --tenant services
```

## Compute Services (Nova)



## Important

Follow the Red Hat documentation's [Chapter 8. Install The Compute Service](#) instructions, but **note the following differences**.

## Controller Node



## Important

Follow the Red Hat documentation's [8.2. Install a Compute Node](#) instructions, but **note the following differences and additions**.

### 1. 8.2.1. Install the Compute Service Packages

Do **not** apply as is.

Instead, install only the following packages:

```
# yum install openstack-nova-api openstack-nova-conductor openstack-nova-scheduler python-cinderclient
```



### Note

The `openstack-nova-compute` package is going to be installed on the Compute Node instead.

#### 2. 8.2.2. Create the Compute Service Database

Apply as is.

#### 3. 8.2.3. Configure the Compute Service Database Connection

Apply as is.

#### 4. 8.2.4. Create the Compute Service Identity Records

Apply as is.

#### 5. 8.2.5. Configure Compute Service Authentication

Apply as is.

#### 6. 8.2.6. Configure the Firewall to Allow Compute Service Traffic

Apply as is.

#### 7. 8.2.7. Configure the Compute Service to Use SSL

Apply as is.

#### 8. 8.2.8. Configure RabbitMQ Message Broker Settings for the Compute Service

Apply as is.

#### 9. 8.2.9. Enable SSL Communication Between the Compute Service and the Message Broker

Apply as is.

#### 10. 8.2.10. Configure Resource Overcommitment

Apply as is.

#### 11. 8.2.11. Reserve Host Resources

Apply as is.

#### 12. 8.2.12. Configure Compute Networking

Apply as is, except the following topics:

##### a. 8.2.12.3. Configure the L2 Agent

Do not apply.

**b. 8.2.12.4. Configure Virtual Interface Plugging**

Configure the generic VIF driver.

**13.8.2.13. Populate the Compute Service Database**

Apply as is.

**14.8.2.14. Launch the Compute Services****a. 1. Starting the Message Bus Service**

Do **not** apply. Only required on the Compute Node.

**b. 2. Starting the Libvirtd Service**

Do **not** apply. Only required on the Compute Node.

**c. 3. Starting the API Service**

Apply as is.

**d. 4. Starting the Scheduler**

Apply as is.

**e. 5. Starting the Conductor**

Apply as is.

**f. 6. Starting the Compute Service**

Do **not** apply. Only required on the Compute Node.

## Compute Node



### Important

Follow the Red Hat documentation's [8.2. Install a Compute Node](#) instructions, but **note the following differences and additions**.

**1. 8.2.1. Install the Compute Service Packages**

Do **not** apply as is.

Instead, install only the following packages:

```
# yum install openstack-nova-compute openstack-utils
```

**2. 8.2.2. Create the Compute Service Database**

Do **not** apply. Has been done on the Controller Node.

**3. 8.2.3. Configure the Compute Service Database Connection**

Apply as is.

**4. 8.2.4. Create the Compute Service Identity Records**

Do **not** apply. Has been done on the Controller Node.

1. **8.2.5. Configure Compute Service Authentication**

Apply as is.

1. **8.2.6. Configure the Firewall to Allow Compute Service Traffic**

Apply as is.

2. **8.2.7. Configure the Compute Service to Use SSL**

Apply as is.

3. **8.2.8. Configure RabbitMQ Message Broker Settings for the Compute Service**

Apply as is.

4. **8.2.9. Enable SSL Communication Between the Compute Service and the Message Broker**

Apply as is.

5. **8.2.10. Configure Resource Overcommitment**

Apply as is.

6. **8.2.11. Reserve Host Resources**

Apply as is.

7. **8.2.12. Configure Compute Networking**

Apply as is, except the following topics:

a. **8.2.12.3. Configure the L2 Agent**

Do **not** apply.

b. **8.2.12.4. Configure Virtual Interface Plugging**

Do **not** apply.

8. **8.2.13. Populate the Compute Service Database**

Do **not** apply. Has been done on the Controller Node.

9. **8.2.14. Launch the Compute Services**

a. **1. Starting the Message Bus Service**

Apply as is.

b. **2. Starting the Libvirtd Service**

Apply as is.

c. **3. Starting the API Service**

Do **not** apply. Only required on the Controller Node.

**d. 4. Starting the Scheduler**

Do **not** apply. Only required on the Controller Node.

**e. 5. Starting the Conductor**

Do **not** apply. Only required on the Controller Node.

**f. 6. Starting the Compute Service**

Apply as is.

**10 Additionally, perform the following steps****a. Configure libvirt**

Edit the `/etc/libvirt/qemu.conf` file to contain the following:

```
user = "root"
group = "root"

cgroup_device_acl = [
    "/dev/null", "/dev/full", "/dev/zero",
    "/dev/random", "/dev/urandom",
    "/dev/ptmx", "/dev/kvm", "/dev/kqemu",
    "/dev/rtc", "/dev/hpet", "/dev/vfio/vfio",
    "/dev/net/tun"
]
```

**b. Restart the libvirt service**

```
# systemctl restart libvirtd.service
```

**c. Install nova-rootwrap network filters**

```
# yum install openstack-nova-network
# systemctl disable openstack-nova-network.service
```

**d. Restart the Compute service**

```
# systemctl restart openstack-nova-compute.service
```

## Networking Services (Neutron)

### Controller Node



#### Important

Follow the Red Hat documentation's [Chapter 7. Install OpenStack Networking](#) instructions, but **note the following differences**.

**1. 7.1. Install the OpenStack Networking Packages**

Do **not** apply as is.

Instead, install the following packages:

```
# yum install openstack-neutron openstack-utils openstack-selinux
python-neutron-plugin-midonet
```

### 2. 7.2.1. Set the OpenStack Networking Plug-in

Do not apply. Instead, perform the following steps:

- a. Edit the `/etc/neutron/neutron.conf` file and configure the following keys in the `[DEFAULT]` section:

```
[DEFAULT]
...
core_plugin = midonet.neutron.plugin_v2.MidonetPluginV2
allow_overlapping_ips = True
```

- b. Create the directory for the MidoNet plugin:

```
mkdir /etc/neutron/plugins/midonet
```

- c. Create the `/etc/neutron/plugins/midonet/midonet.ini` file and edit it to contain the following:

```
[DATABASE]
sql_connection = mysql://neutron:NEUTRON_DBPASS@controller/neutron

[MIDONET]
# MidoNet API URL
midonet_uri = http://controller:8181/midonet-api
# MidoNet administrative user in Keystone
username = midonet
password = MIDONET_PASS
# MidoNet administrative user's tenant
project_id = services
```

- d. Create a symbolic link to direct Neutron to the MidoNet configuration:

```
# ln -s /etc/neutron/plugins/midonet/midonet.ini /etc/neutron/plugin.ini
```

### 3. 7.2.2. Create the OpenStack Networking Database

Do not apply.

Instead, create the database as follows:

```
$ mysql -u root -p
CREATE DATABASE neutron character set utf8;
GRANT ALL PRIVILEGES ON neutron.* TO 'neutron'@'localhost' IDENTIFIED BY
'NEUTRON_DBPASS';
GRANT ALL PRIVILEGES ON neutron.* TO 'neutron'@'%' IDENTIFIED BY
'NEUTRON_DBPASS';
FLUSH PRIVILEGES;
quit
```

Afterwards, run the `neutron-db-manage` command:

```
# neutron-db-manage \
--config-file /usr/share/neutron/neutron-dist.conf \
--config-file /etc/neutron/neutron.conf \
--config-file /etc/neutron/plugin.ini \
upgrade head
```

Followed by the `midonet-db-manage` command:

```
# midonet-db-manage upgrade head
```

### 4. 7.2.3. Configure the OpenStack Networking Database Connection

Apply as is.

#### 5. 7.2.4. Create the OpenStack Networking Identity Records

Apply as is.

#### 6. 7.2.5. Configure OpenStack Networking Authentication

Apply as is.

#### 7. 7.2.6. Configure the Firewall to Allow OpenStack Networking Traffic

Apply as is.

#### 8. 7.2.7. Configure RabbitMQ Message Broker Settings for OpenStack Networking

Apply as is.

#### 9. 7.2.8. Enable SSL Communication Between OpenStack Networking and the Message Broker

Apply as is.

#### 10. 7.2.9. Configure OpenStack Networking to Communicate with the Compute Service

Apply as is.

#### 11. Configure Load-Balancer-as-a-Service (LBaaS)

Additionally to the Red Hat Installation Guide, configure Load-Balancer-as-a-Service (LBaaS) as described in [the section called "Configure Load-Balancer-as-a-Service \(LBaaS\)" \[13\]](#).

#### 12. 7.2.10. Launch OpenStack Networking

Apply as is.

#### 13. 7.3. Configure the DHCP Agent

Do not apply.

#### 14. 7.4. Create an External Network

Do not apply.

Instead, create the Neutron networks after the OpenStack and MidoNet installation is completed.

Any networks that are created before the MidoNet plug-in is active will not be visible to MidoNet.

#### 15. 7.5. Configure the Plug-in Agent

Do not apply.

#### 16. 7.6. Configure the L3 Agent

Do not apply.



## Configure Load-Balancer-as-a-Service (LBaaS)



### Important

Follow the OpenStack documentation's [Configure Load-Balancer-as-a-Service \(LBaaS\)](#) instructions, but **note the following differences**.

#### 1. Install the agent

Do not apply.

Instead install the package `python-neutron-lbaas`

```
# yum install python-neutron-lbaas
```

#### 2. Enable the HAProxy plug-in by using the `service_provider` option in the `/etc/neutron/neutron.conf` file:

Do not apply.

Instead, set `service_provider` as follows:

```
[service_providers]
service_provider = LOADBALANCER:Midonet:midonet.neutron.services.
loadbalancer.driver.MidonetLoadbalancerDriver:default
```

#### 3. Enable the load-balancing plug-in by using the `service_plugins` option in the `/etc/neutron/neutron.conf` file:

Apply as is.

#### 4. Enable the HAProxy load balancer in the `/etc/neutron/lbaas_agent.ini` file:

Do not apply.

#### 5. Select the required driver in the `/etc/neutron/lbaas_agent.ini` file:

Do not apply.

#### 6. Create the required tables in the database:

Do not apply.

#### 7. Apply the settings by restarting the `neutron-server` and `neutron-lbaas-agent` services.

Do not apply.

#### 8. Enable load balancing in the Project section of the dashboard.

Apply as is.

#### 9. To finalize installation

Finalize the installation as described in [Neutron Controller Node Installation](#).

## 4. MidoNet Installation

### Table of Contents

NSDB Nodes .....	14
Controller Node .....	17
Midolman Installation .....	19
MidoNet Host Registration .....	20

## NSDB Nodes

### ZooKeeper Installation

#### 1. Install ZooKeeper packages

```
# yum install java-1.7.0-openjdk-headless
# yum install zookeeper zkdump nmap-ncat
```

#### 2. Configure ZooKeeper

##### a. Common Configuration

Edit the `/etc/zookeeper/zoo.cfg` file to contain the following:

```
server.1=nsdb1:2888:3888
server.2=nsdb2:2888:3888
server.3=nsdb3:2888:3888
```

Create data directory:

```
# mkdir /var/lib/zookeeper/data
# chown zookeeper:zookeeper /var/lib/zookeeper/data
```



#### Important

For production deployments it is recommended to configure the storage of snapshots in a different disk than the commit log. This can be set by changing the parameter `dataDir` in `zoo.cfg` to a different disk.

##### b. Node-specific Configuration

###### i. NSDB Node 1

Create the `/var/lib/zookeeper/data/myid` file and edit it to contain the host's ID:

```
# echo 1 > /var/lib/zookeeper/data/myid
```

###### ii. NSDB Node 2

Create the `/var/lib/zookeeper/data/myid` file and edit it to contain the host's ID:

```
# echo 2 > /var/lib/zookeeper/data/myid
```

### iii. NSDB Node 3

Create the `/var/lib/zookeeper/data/myid` file and edit it to contain the host's ID:

```
# echo 3 > /var/lib/zookeeper/data/myid
```

### 3. Create Java Symlink

```
# mkdir -p /usr/java/default/bin/  
# ln -s /usr/lib/jvm/jre-1.7.0-openjdk/bin/java /usr/java/default/bin/  
java
```

### 4. Enable and start ZooKeeper

```
# systemctl enable zookeeper.service  
# systemctl start zookeeper.service
```

### 5. Verify ZooKeeper Operation

After installation of all nodes has been completed, verify that ZooKeeper is operating properly.

A basic check can be done by executing the `ruok` (Are you ok?) command on all nodes. This will reply with `imok` (I am ok.) if the server is running in a non-error state:

```
$ echo ruok | nc 127.0.0.1 2181  
imok
```

More detailed information can be requested with the `stat` command, which lists statistics about performance and connected clients:

```
$ echo stat | nc 127.0.0.1 2181  
Zookeeper version: 3.4.5--1, built on 06/10/2013 17:26 GMT  
Clients:  
 /127.0.0.1:34768[0](queued=0,recved=1,sent=0)  
 /192.0.2.1:49703[1](queued=0,recved=1053,sent=1053)  
  
Latency min/avg/max: 0/4/255  
Received: 1055  
Sent: 1054  
Connections: 2  
Outstanding: 0  
Zxid: 0x260000013d  
Mode: follower  
Node count: 3647
```

## Cassandra Installation

### 1. Install Cassandra packages

```
# yum install java-1.7.0-openjdk-headless  
# yum install dsc20
```

### 2. Configure Cassandra

#### a. Common Configuration

Edit the `/etc/cassandra/conf/cassandra.yaml` file to contain the following:

```
# The name of the cluster.
```

```
cluster_name: 'midonet'

...

# Addresses of hosts that are deemed contact points.
seed_provider:
  - class_name: org.apache.cassandra.locator.SimpleSeedProvider
    parameters:
      - seeds: "nsdb1,nsdb2,nsdb3"
```

## b. Node-specific Configuration

### i. NSDB Node 1

Edit the `/etc/cassandra/conf/cassandra.yaml` file to contain the following:

```
# Address to bind to and tell other Cassandra nodes to connect to.
listen_address: nsdb1

...

# The address to bind the Thrift RPC service.
rpc_address: nsdb1
```

### ii. NSDB Node 2

Edit the `/etc/cassandra/conf/cassandra.yaml` file to contain the following:

```
# Address to bind to and tell other Cassandra nodes to connect to.
listen_address: nsdb2

...

# The address to bind the Thrift RPC service.
rpc_address: nsdb2
```

### iii. NSDB Node 3

Edit the `/etc/cassandra/conf/cassandra.yaml` file to contain the following:

```
# Address to bind to and tell other Cassandra nodes to connect to.
listen_address: nsdb3

...

# The address to bind the Thrift RPC service.
rpc_address: nsdb3
```

## 3. Edit the service's init script

On installation, the `/var/run/cassandra` directory is created, but because it is located on a temporary file system it will be lost after system reboot. As a result it is not possible to stop or restart the Cassandra service anymore.

To avoid this, edit the `/etc/init.d/cassandra` file to create the directory on service start:

```
[...]
case "$1" in
  start)
```

```
# Cassandra startup
echo -n "Starting Cassandra: "
mkdir -p /var/run/cassandra
chown cassandra:cassandra /var/run/cassandra
su $CASSANDRA_OWNRR -c "$CASSANDRA_PROG -p $pid_file" > $log_file

2>&1

retval=$?

[...]
```

#### 4. Enable and start Cassandra

```
# systemctl enable cassandra.service
# systemctl start cassandra.service
```

## 5. Verify Cassandra Operation

After installation of all nodes has been completed, verify that Cassandra is operating properly.



## Important

If Cassandra fails to start and prints a "buffer overflow" error message in its log file, you may try associating `127.0.0.1` with the hostname in `etc/hosts` (so that `hostname -i` will show `127.0.0.1`). This may solve the Cassandra start problem.

A basic check can be done by executing the `nodetool status` command. This will reply with `UN` (Up / Normal) in the first column if the servers are running in a non-error state:

```
$ nodetool -host 127.0.0.1 status
[...]
```

Status=Up/Down					
/ State=Normal/Leaving/Joining/Moving					
--	Address	Load	Tokens	Owns	Host ID
	Rack				
UN	192.0.2.1	123.45 KB	256	33.3%	
	11111111-2222-3333-4444-555555555555				rack1
UN	192.0.2.2	234.56 KB	256	33.3%	
	22222222-3333-4444-5555-666666666666				rack1
UN	192.0.2.3	345.67 KB	256	33.4%	
	33333333-4444-5555-6666-777777777777				rack1

## Controller Node

## MidoNet Cluster Installation

## 1. Install MidoNet Cluster package

```
# yum install midonet-cluster
```

## 2. Set up mn-conf

Edit `/etc/midonet/midonet.conf` to point `mn-conf` to the ZooKeeper cluster:

```
[zookeeper]
zookeeper_hosts = nsdb1:2181,nsdb2:2181,nsdb3:2181
```

This step needs to happen only once, it will set up access to the NSDB for the MidoNet Cluster and Agent nodes.

Run the following command to set the cloud-wide values for the ZooKeeper and Cassandra server addresses:

```
$ cat << EOF | mn-conf set -t default
zookeeper {
    zookeeper_hosts = "nsdb1:2181,nsdb2:2181,nsdb3:2181"
}

cassandra {
    servers = "nsdb1,nsdb2,nsdb3"
}
EOF
```

Run the following command to set the Cassandra replication factor:

```
$ echo "cassandra.replication_factor : 3" | mn-conf set -t default
```

#### 4. Configure Keystone access

This step needs to happen only once, it will set up access to Keystone for the MidoNet Cluster node(s).

This step will configure the local Midonet Cluster node to be able to use Keystone.

```
$ cat << EOF | mn-conf set -t default
cluster.auth {
    provider_class =
    "org.midonet.cluster.auth.keystone.v2_0.KeystoneService"
    admin_role = "admin"
    keystone.tenant_name = "admin"
    keystone.admin_token = "ADMIN_TOKEN"
    keystone.host = controller
    keystone.port = 35357
}
EOF
```

#### 5. Start the MidoNet Cluster

```
# systemctl start midonet-cluster.service
```

## MidoNet CLI Installation

#### 1. Install MidoNet CLI package

```
# yum install python-midonetclient
```

#### 2. Configure MidoNet CLI

Create the `~/.midonetr` file and edit it to contain the following:

```
[cli]
api_url = http://controller:8181/midonet-api
username = admin
password = ADMIN_PASS
project_id = admin
```

# Midolman Installation

The *MidoNet Agent (Midolman)* has to be installed on all nodes where traffic enters or leaves the virtual topology, in this guide this are the **gateway1**, **gateway2** and **compute1** nodes.

## 1. Install Midolman package

```
# yum install java-1.8.0-openjdk-headless
# yum install midolman
```

## 2. Configure resource usage

Run these steps on **each agent host** in order to configure resource usage.



### Important

For production environments the **large** templates are strongly recommended.

### a. Midolman resource template

Run the following command to configure the Midolman resource template:

```
$ mn-conf template-set -h local -t TEMPLATE_NAME
```

Replace **TEMPLATE\_NAME** with one of the following templates:

```
agent-compute-large
agent-compute-medium
agent-gateway-large
agent-gateway-medium
default
```

### b. Java Virtual Machine (JVM) resource template

Replace the default `/etc/midolman/midolman-env.sh` file with one of the below to configure the JVM resource template:

```
/etc/midolman/midolman-env.sh.compute.large
/etc/midolman/midolman-env.sh.compute.medium
/etc/midolman/midolman-env.sh.gateway.large
/etc/midolman/midolman-env.sh.gateway.medium
```

## 3. Configure MidoNet Metadata Proxy for all agents

This step needs to happen only once, it will set up MidoNet Metadata Proxy for all MidoNet Agent nodes.

Run the following commands to set the cloud-wide values for the MidoNet Metadata Proxy:

```
$ echo "agent.openstack.metadata.nova_metadata_url : \"http://
/nova_metadata_host:nova_metadata_port\"" | mn-conf set
$ echo "agent.openstack.metadata.shared_secret : shared_secret" | mn-
conf set
$ echo "agent.openstack.metadata.enabled : true" | mn-conf set
```

**nova\_metadata\_host**, **nova\_metadata\_port**, and **shared\_secret** should be replaced with appropriate values. They need to match with the corresponding Nova Metadata API configuration.

**nova\_metadata\_host** and **nova\_metadata\_port** specify the address on which Nova accepts Metadata API requests. **shared\_secret** has to be the same as specified by the "metadata\_proxy\_shared\_secret" field in the "neutron" section of nova.conf.

The Nova side of the configuration for the metadata service is same as when using Neutron Metadata Proxy. See the OpenStack documentation for details:

[Cloud Administrator Guide: Configure Metadata](#)

#### 4. Start Midolman

```
# systemctl start midolman.service
```

## MidoNet Host Registration

#### 1. Launch MidoNet CLI

```
$ midonet-cli  
midonet>
```

#### 2. Create tunnel zone

MidoNet supports the Virtual Extensible LAN (VXLAN) and Generic Routing Encapsulation (GRE) protocols to communicate to other hosts within a tunnel zone.

To use the VXLAN protocol, create the tunnel zone with type 'vxlan':

```
midonet> tunnel-zone create name tz type vxlan  
tzone0
```

To use the GRE protocol, create the tunnel zone with type 'gre':

```
midonet> tunnel-zone create name tz type gre  
tzone0
```



### Important

Make sure to allow GRE/VXLAN traffic for all hosts that belong to the tunnel zone. For VXLAN MidoNet uses UDP port 6677 as default.

#### 1. Add hosts to tunnel zone

```
midonet> list tunnel-zone  
tzone tzone0 name tz type vxlan  
  
midonet> list host  
host host0 name controller alive true  
host host1 name gateway1 alive true  
host host2 name gateway2 alive true  
host host3 name compute1 alive true  
  
midonet> tunnel-zone tzone0 add member host host0  
address ip_address_of_host0  
zone tzone0 host host0 address ip_address_of_host0  
  
midonet> tunnel-zone tzone0 add member host host1  
address ip_address_of_host1  
zone tzone0 host host1 address ip_address_of_host1
```



```
midonet> tunnel-zone tzone0 add member host host2
        address ip_address_of_host2
zone tzone0 host host2 address ip_address_of_host2

midonet> tunnel-zone tzone0 add member host host3
        address ip_address_of_host3
zone tzone0 host host3 address ip_address_of_host3
```

MidoNet utilizes the Border Gateway Protocol (BGP) for external connectivity.

For production deployments it is strongly recommended to use BGP due to it's scalability and redundancy.

For demo or POC environments, alternatively static routing can be used. See the Operations Guide for details.

The following instructions assume below sample environment:

- Follow these steps to configure the BGP uplinks.

### 1. Determine the Keystone admin tenant ID

Use the `keystone` command to determine the Keystone `admin` tenant's ID:

```
$ keystone tenant-list
```

id	name	enabled
12345678901234567890123456789012	admin	True

## 2. Launch the MidoNet CLI and find the MidoNet Provider Router

```
$ midonet-cli
midonet-cli>
```

Because the MidoNet Provider Router is not associated with a tenant, the active tenant has to be cleared (`cleart`) first.

```
midonet-cli> clear

midonet-cli> router list
router router0 name MidoNet Provider Router state up
router router1 name Tenant Router state up infiltrer chain0 outfilter
chain1
```

### 3. Load the admin tenant

#### 4. Create virtual ports for the BGP sessions

In this example the created ports are **port0** and **port1**.

## 6. Add routes to the remote BGP peers

## 7. Advertise BGP routes

In order to provide external connectivity for hosted virtual machines, the floating IP network has to be advertised to the BGP peers.

```
midonet> router router0 add bgp-network net 192.0.2.0/24
router0:net0

midonet> router router0 list bgp-network
net net0 net 192.0.2.0/24
```

## 8. Bind virtual ports to physical network interfaces

Bind the MidoNet Provider Router's virtual ports to the physical network interfaces on the Gateway Nodes.



### Important

Ensure that physical interfaces are in state UP and do not have an IP address assigned.

#### a. List the MidoNet hosts and find the Gateway Nodes:

```
midonet> host list
host host0 name gateway1 alive true
host host1 name gateway2 alive true
[...]
```

In this example the hosts are **host0** and **host1**.

#### b. List the Gateway Nodes' physical interfaces:

```
midonet> host host0 list interface
[...]
iface eth1 host_id host0 status 3 addresses [] mac 01:02:03:04:05:06
mtu 1500 type Physical endpoint PHYSICAL
[...]

midonet> host host1 list interface
[...]
iface eth1 host_id host0 status 3 addresses [] mac 06:05:04:03:02:01
mtu 1500 type Physical endpoint PHYSICAL
[...]
```

#### c. Bind the physical host interfaces to the MidoNet Provider Router's virtual ports:

```
midonet> host host0 add binding port router0:port0 interface eth1
host host0 interface eth1 port router0:port0

midonet> host host1 add binding port router0:port1 interface eth1
host host1 interface eth1 port router0:port1
```

#### d. Configure a stateful port group:

```
midonet-cli> port-group create name uplink-spg stateful true
pgroup0
```

#### e. Add the ports to the port group:

```
midonet> port-group pgroup0 add member port router0:port0
port-group pgroup0 port router0:port0

midonet> port-group pgroup0 add member port router0:port1
port-group pgroup0 port router0:port1
```

```
midonet> port-group pgroup0 list member
port-group pgroup0 port router0:port0
port-group pgroup0 port router0:port1
```

## 6. Further Steps

MidoNet installation and integration into OpenStack is completed.

You can now continue with the creation of initial networks in Neutron.



### Note

Consult the **Operations Guide** for further instructions on operating MidoNet.