



MidoNet Quick Start Guide

for SLES 12 / Liberty (SOC)

5.1-SNAPSHOT (2016-05-12 02:14 UTC)

DRAFT

MidoNet Quick Start Guide for SLES 12 / Liberty (SOC)

5.1-SNAPSHOT (2016-05-12 02:14 UTC)

Copyright © 2016 Midokura SARL All rights reserved.

MidoNet is a network virtualization software for Infrastructure-as-a-Service (IaaS) clouds.

It decouples your IaaS cloud from your network hardware, creating an intelligent software abstraction layer between your end hosts and your physical network.

This guide walks through the minimum installation and configuration steps necessary to use MidoNet with OpenStack.



Caution

This document is a DRAFT. It may be MISSING relevant information or contain UNTESTED information. Use it at your own risk.



Note

Please consult the [MidoNet Mailing Lists or Chat](#) if you need assistance.

Licensed under the Apache License, Version 2.0 (the "License"); you may not use this file except in compliance with the License. You may obtain a copy of the License at

<http://www.apache.org/licenses/LICENSE-2.0>

Unless required by applicable law or agreed to in writing, software distributed under the License is distributed on an "AS IS" BASIS, WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or implied. See the License for the specific language governing permissions and limitations under the License.

Table of Contents

Preface	iv
Conventions	iv
1. Architecture	1
Hosts and Services	2
2. Basic Environment Configuration	4
Networking Configuration	4
SuSEfirewall2 Configuration	4
Repository Configuration	4
3. OpenStack Installation	5
Identity Service (Keystone)	5
Compute Services (Nova)	6
Networking Services (Neutron)	8
4. MidoNet Installation	13
NSDB Nodes	13
Controller Node	16
Midolman Installation	18
MidoNet Host Registration	19
5. Initial Network Configuration	21
6. Edge Router Setup	22
7. BGP Uplink Configuration	24
8. Further Steps	26

Preface

Conventions

The MidoNet documentation uses several typesetting conventions.

Notices

Notices take these forms:



Note

A handy tip or reminder.



Important

Something you must be aware of before proceeding.



Warning

Critical information about the risk of data loss or security issues.

Command prompts

\$ prompt

Any user, including the root user, can run commands that are prefixed with the \$ prompt.

prompt

The root user must run commands that are prefixed with the # prompt. You can also prefix these commands with the **sudo** command, if available, to run them.

Table of Contents

This guide assumes the following example system architecture.

- Controller Node (**controller**)

- Compute Node (**compute1**)

- Gateway Node 1 (**gateway1**)

- Gateway Node 2 (**gateway2**)

Three or more hosts are being used for the MidoNet Network State Database (NSDB) cluster which utilizes [ZooKeeper](#) and [Cassandra](#) to store virtual network topology and connection state information:

- NSDB Node 1 (**nsdb1**)
- NSDB Node 2 (**nsdb2**)
- NSDB Node 3 (**nsdb3**)



Ideally, both the ZooKeeper transaction log and Cassandra data files need their own dedicated disks, with additional disks for other services on the host. However, for small POCs and small deployments, it is ok to share the Cassandra disk with other services and just leave the ZooKeeper transaction log on its own.

The *MidoNet Agent (Midolman)* has to be installed on all nodes where traffic enters or leaves the virtual topology. In this guide this are the **gateway1**, **gateway2** and **compute1** hosts.

The *Midonet Cluster* can be installed on a separate host, but this guide assumes it to be installed on the **controller** host.

The *MidoNet Command Line Interface (CLI)* can be installed on any host that has connectivity to the MidoNet Cluster. This guide assumes it to be installed on the **controller** host.

The *Midonet Neutron Plugin* replaces the ML2 Plugin and has to be installed on the **controller**.

Hosts and Services

Controller Node (controller)

- General
 - Database (MariaDB)
 - Message Broker (RabbitMQ)
- OpenStack
 - Identity Service (Keystone)
 - Image Service (Glance)
 - Compute (Nova)
 - Networking (Neutron)
 - Neutron Server
 - Dashboard (Horizon)
- MidoNet
 - Cluster
 - CLI
 - Neutron Plugin

Compute Node (compute1)

- OpenStack
 - Compute (Nova)
 - Networking (Neutron)
- MidoNet
 - Agent (Midolman)

NSDB Nodes (nsdb1, nsdb2, nsdb3)

- Network State Database (NSDB)
 - Network Topology (ZooKeeper)
 - Network State Information (Cassandra)

Gateway Nodes (gateway1, gateway2)

- BGP Daemon (Quagga)

- MidoNet
 - Agent (Midolman)

2. Basic Environment Configuration

Table of Contents

Networking Configuration	4
SuSEfirewall2 Configuration	4
Repository Configuration	4

Networking Configuration



Important

All hostnames must be resolvable, either via DNS or locally.

This guide assumes that you follow the instructions in [Network](#) of the SUSE OpenStack Cloud Documentation.

SuSEfirewall2 Configuration



Note

SuSEfirewall2 is a stateful network packet filter also known as firewall.

Disable it by executing the following command:

```
# /sbin/SuSEfirewall2 off
```

Repository Configuration

Configure necessary software repositories and update installed packages.

1. Enable MidoNet repositories

Use the following commands to add the MidoNet repositories:

```
# zypper addrepo -f http://builds.midonet.org/midonet-5/stable/sles12/  
midonet
```

```
# zypper addrepo -f http://builds.midonet.org/openstack-liberty/stable/  
sles12/ midonet-plugin-liberty
```

Download and import the repositories' key:

```
# wget https://builds.midonet.org/midorepo.key  
# rpm --import midorepo.key
```

2. Install available updates

```
# zypper update
```

3. If necessary, reboot the system

```
# reboot
```


3. OpenStack Installation

Table of Contents

Identity Service (Keystone)	5
Compute Services (Nova)	6
Networking Services (Neutron)	8



Important

Follow the [Suse Cloud 6 Deployment Guide](#) , but **note the following differences.**

Identity Service (Keystone)



Important

Follow the SUSE OpenStack Cloud documentation's [Deploying Keystone](#) instructions, but **note the following additions**.

1. Verify Admin Token Authentication

The MidoNet API uses the Keystone admin token for authentication purposes, therefore `admin_token_auth` has to be kept in the corresponding Keystone pipeline configuration sections.

Ensure that `admin_token_auth` is included in the following pipelines in `/usr/share/keystone/keystone-dist-paste.ini`:

```
[pipeline:public_api]
pipeline = [...] admin_token_auth [...]

[pipeline:admin_api]
pipeline = [...] admin_token_auth [...]

[pipeline:api_v3]
pipeline = [...] admin_token_auth [...]
```

2. Create MidoNet API Service

As Keystone admin, execute the following command:

```
$ openstack service create --name midonet --description "MidoNet API Service" midonet
```

3. Create MidoNet Administrative User

As Keystone admin, execute the following commands:

```
$ openstack user create --domain default --password-prompt midonet
$ openstack role add --project service --user midonet admin
```

Compute Services (Nova)



Important

Follow the SUSE OpenStack Cloud documentation's [Deploying Nova](#) instructions, but **note the following differences**.

Controller Node



Note

Follow the SUSE OpenStack Cloud documentation's [Deploying the OpenStack Services](#) instructions as is.

Compute Node



Important

Follow the OpenStack documentation's [Deploying Nova](#) instructions, but **note the following additions**.

1. Configure libvirt

Edit the `/etc/libvirt/qemu.conf` file to contain the following:

```
user = "root"
group = "root"

cgroup_device_acl = [
    "/dev/null", "/dev/full", "/dev/zero",
    "/dev/random", "/dev/urandom",
    "/dev/ptmx", "/dev/kvm", "/dev/kqemu",
    "/dev/rtc", "/dev/hpet", "/dev/vfio/vfio",
    "/dev/net/tun"
]
```

2. Restart the libvirt service

```
# service libvirtd restart
```

3. Edit the nova-rootwrap network filters `/etc/nova/rootwrap.d/network.filters` file to contain the following

```
# nova-rootwrap command filters for network nodes
# This file should be owned by (and only-writeable by) the root user

[Filters]
# nova/virt/libvirt/vif.py: 'ip', 'tuntap', 'add', dev, 'mode', 'tap'
# nova/virt/libvirt/vif.py: 'ip', 'link', 'set', dev, 'up'
# nova/virt/libvirt/vif.py: 'ip', 'link', 'delete', dev
# nova/network/linux_net.py: 'ip', 'addr', 'add', str(floating_ip)+'/'
32'i..
# nova/network/linux_net.py: 'ip', 'addr', 'del', str(floating_ip)+'/'
32'..
# nova/network/linux_net.py: 'ip', 'addr', 'add', '169.254.169.254/32',..
.
# nova/network/linux_net.py: 'ip', 'addr', 'show', 'dev', dev, 'scope',..
.
# nova/network/linux_net.py: 'ip', 'addr', 'del/add', ip_params, dev)
# nova/network/linux_net.py: 'ip', 'addr', 'del', params, fields[-1]
```

```
# nova/network/linux_net.py: 'ip', 'addr', 'add', params, bridge
# nova/network/linux_net.py: 'ip', '-f', 'inet6', 'addr', 'change', ..
# nova/network/linux_net.py: 'ip', 'link', 'set', 'dev', dev, 'promisc',
..
# nova/network/linux_net.py: 'ip', 'link', 'add', 'link', bridge_if ...
# nova/network/linux_net.py: 'ip', 'link', 'set', interface, address,..
# nova/network/linux_net.py: 'ip', 'link', 'set', interface, 'up'
# nova/network/linux_net.py: 'ip', 'link', 'set', bridge, 'up'
# nova/network/linux_net.py: 'ip', 'addr', 'show', 'dev', interface, ..
# nova/network/linux_net.py: 'ip', 'link', 'set', dev, address, ..
# nova/network/linux_net.py: 'ip', 'link', 'set', dev, 'up'
# nova/network/linux_net.py: 'ip', 'route', 'add', ..
# nova/network/linux_net.py: 'ip', 'route', 'del', .
# nova/network/linux_net.py: 'ip', 'route', 'show', 'dev', dev
ip: CommandFilter, ip, root

# nova/virt/libvirt/vif.py: 'ovs-vsctl', ...
# nova/virt/libvirt/vif.py: 'ovs-vsctl', 'del-port', ...
# nova/network/linux_net.py: 'ovs-vsctl', ....
ovs-vsctl: CommandFilter, ovs-vsctl, root

# nova/network/linux_net.py: 'ovs-ofctl', ....
ovs-ofctl: CommandFilter, ovs-ofctl, root

# nova/virt/libvirt/vif.py: 'ivs-ctl', ...
# nova/virt/libvirt/vif.py: 'ivs-ctl', 'del-port', ...
# nova/network/linux_net.py: 'ivs-ctl', ....
ivs-ctl: CommandFilter, ivs-ctl, root

# nova/virt/libvirt/vif.py: 'ifc_ctl', ...
ifc_ctl: CommandFilter, /opt/pg/bin/ifc_ctl, root

# nova/virt/libvirt/vif.py: 'ebrctl', ...
ebrctl: CommandFilter, ebrctl, root

# nova/virt/libvirt/vif.py: 'mm-ctl', ...
mm-ctl: CommandFilter, mm-ctl, root

# nova/network/linux_net.py: 'ebtables', '-D' ...
# nova/network/linux_net.py: 'ebtables', '-I' ...
ebtables: CommandFilter, ebtables, root
ebtables_usr: CommandFilter, ebtables, root

# nova/network/linux_net.py: 'ip[6]tables-save' % (cmd, '-t', ...
iptables-save: CommandFilter, iptables-save, root
ip6tables-save: CommandFilter, ip6tables-save, root

# nova/network/linux_net.py: 'ip[6]tables-restore' % (cmd,)
iptables-restore: CommandFilter, iptables-restore, root
ip6tables-restore: CommandFilter, ip6tables-restore, root

# nova/network/linux_net.py: 'arping', '-U', floating_ip, '-A', '-I', ..
.
# nova/network/linux_net.py: 'arping', '-U', network_ref['dhcp_server'],
..
arping: CommandFilter, arping, root

# nova/network/linux_net.py: 'dhcp_release', dev, address, mac_address
dhcp_release: CommandFilter, dhcp_release, root

# nova/network/linux_net.py: 'kill', '-9', pid
# nova/network/linux_net.py: 'kill', '-HUP', pid
kill_dnsmasq: KillFilter, root, /usr/sbin/dnsmasq, -9, -HUP

# nova/network/linux_net.py: 'kill', pid
```

```
kill_radvd: KillFilter, root, /usr/sbin/radvd

# nova/network/linux_net.py: dnsmasq call
dnsmasq: EnvFilter, env, root, CONFIG_FILE=, NETWORK_ID=, dnsmasq

# nova/network/linux_net.py: 'radvd', '-C', '%s' % _ra_file(dev, 'conf'.
.
radvd: CommandFilter, radvd, root

# nova/network/linux_net.py: 'brctl', 'addbr', bridge
# nova/network/linux_net.py: 'brctl', 'setfd', bridge, 0
# nova/network/linux_net.py: 'brctl', 'stp', bridge, 'off'
# nova/network/linux_net.py: 'brctl', 'addif', bridge, interface
brctl: CommandFilter, brctl, root

# nova/network/linux_net.py: 'sysctl', ....
sysctl: CommandFilter, sysctl, root

# nova/network/linux_net.py: 'conntrack'
conntrack: CommandFilter, conntrack, root
```

4. Restart the Compute service

```
# service openstack-nova-compute restart
```

Networking Services (Neutron)



Important

Follow the SUSE OpenStack Cloud documentation's [Deploying Neutron](#) instructions, but **note the following differences**.

Controller Node



Important

Follow the OpenStack documentation's [Deploying Nova](#) instructions, but **note the following differences and additions**.

1. Prerequisites

Apply as is.

2. Configure networking options

Do **not** apply.

a. Instead, install the following packages:

```
# zypper install openstack-neutron python-networking-midonet python-
neutronclient
```

b. Configure the server component:

Edit the `/etc/neutron/neutron.conf` file and configure the following keys:

```
[DEFAULT]
...
core_plugin = midonet.neutron.plugin_v2.MidonetPluginV2
```

```
service_plugins
= midonet.neutron.services.l3.l3_midoNet.MidoNetL3ServicePlugin
...
dhcp_agent_notification = False
...
allow_overlapping_ips = True
...
rpc_backend = rabbit
...
auth_strategy = keystone
...
notify_nova_on_port_status_changes = True
notify_nova_on_port_data_changes = True
nova_url = http://controller:8774/v2

[database]
...
connection = postgresql://neutron:NEUTRON_DBPASS@controller/neutron

[oslo_messaging_rabbit]
...
rabbit_host = controller
rabbit_userid = openstack
rabbit_password = RABBIT_PASS

[keystone_authtoken]
...
auth_uri = http://controller:5000
auth_url = http://controller:35357
auth_plugin = password
project_domain_id = default
user_domain_id = default
project_name = service
username = neutron
password = NEUTRON_PASS

[nova]
...
auth_url = http://controller:35357
auth_plugin = password
project_domain_id = default
user_domain_id = default
region_name = RegionOne
project_name = service
username = nova
password = NOVA_PASS

[oslo_concurrency]
...
lock_path = /var/lib/neutron/tmp
```



Note

When using multiple service plugins, separate them with commas:

```
[DEFAULT]
service_plugins = foo,bar,midonet.neutron.services.l3.
l3_midoNet.MidoNetL3ServicePlugin
```

3. Configure the MidoNet plug-in

- a. Create the directory for the MidoNet plugin:

```
mkdir /etc/neutron/plugins/midoNet
```

- b. Create the `/etc/neutron/plugins/midonet/midonet.ini` file and edit it to contain the following:

```
[DATABASE]
sql_connection = postgresql://neutron:NEUTRON_DBPASS@controller/
neutron

[MIDONET]
# MidoNet API URL
midonet_uri = http://controller:8080/midonet-api
# MidoNet administrative user in Keystone
username = midonet
password = MIDONET_PASS
# MidoNet administrative user's tenant
project_id = service
```

- c. Create a symbolic link to direct Neutron to the MidoNet configuration:

```
# ln -s /etc/neutron/plugins/midonet/midonet.ini /etc/neutron/plugin.
ini
```

4. Configure the metadata agent

Do not apply.

5. To configure Compute to use Networking

Apply as is.

6. Configure Load-Balancer-as-a-Service (LBaaS)

Additionally to the OpenStack Installation Guide, configure Load-Balancer-as-a-Service (LBaaS) as described in [the section called "Configure Load-Balancer-as-a-Service \(LBaaS\)" \[10\]](#).

7. To finalize installation

Do not apply.

Instead, perform the following steps.

- a. Populate the database:

```
# su -s /bin/sh -c "neutron-db-manage --config-file /etc/neutron/
neutron.conf --config-file /etc/neutron/plugins/midonet/midonet.ini
upgrade head" neutron
# su -s /bin/sh -c "neutron-db-manage --subproject networking-midonet
upgrade head" neutron
```

- b. Restart the Compute service:

```
# service openstack-nova-api restart
```

- c. Start the Networking service and configure it to start when the system boots:

```
# service neutron-server start
```

Configure Load-Balancer-as-a-Service (LBaaS)

1. Install Neutron Load-Balancing-as-a-Service

```
# zypper install python-neutron-lbaas
```

2. Enable the MidoNet driver

Enable the MidoNet driver by using the `service_provider` option in the `/etc/neutron/neutron.conf` file:

```
[service_providers]
service_provider = LOADBALANCER:Midonet:midonet.neutron.services.
loadbalancer.driver.MidonetLoadbalancerDriver:default
```

3. Enable the LBaaS plug-in

Enable the LBaaS plug-in by using the `service_plugins` option in the `/etc/neutron/neutron.conf` file:

```
service_plugins = lbaas
```



Note

When using multiple service plugins, separate them with commas:

```
[DEFAULT]
service_plugins = foo,bar,lbaas
```

4. Enable load balancing in the dashboard

Change the `enable_lb` option to `True` in the `/srv/www/openstack-dashboard/openstack_dashboard/local/local_settings.py` file:

```
OPENSTACK_NEUTRON_NETWORK = {
    'enable_lb': True,
    ...
}
```

5. To finalize installation

Finalize the installation as described in [Neutron Controller Node Installation](#).

Configure FireWall-as-a-Service (FWaaS)

1. Install Neutron FireWall-as-a-Service

```
# zypper install python-neutron-fwaas
```

2. Enable the MidoNet FWaaS plug-in

Enable the MidoNet FWaaS plug-in by using the `service_plugins` option in the `/etc/neutron/neutron.conf` file:

```
service_plugins = midonet.neutron.services.firewall.plugin.
MidonetFirewallPlugin
```



Note

When using multiple service plugins, separate them with commas:

```
[DEFAULT]
service_plugins = foo,bar,midonet.neutron.services.firewall.
plugin.MidonetFirewallPlugin
```

3. Enable firewall in the dashboard 11

Change the `enable_firewall` option to `True` in the `/etc/openstack-dashboard/local_settings.py` file:

```
OPENSTACK_NEUTRON_NETWORK = {  
    'enable_firewall': True,  
    ...  
}
```

4. To finalize installation

Finalize the installation as described in [Neutron Controller Node Installation](#).

Compute Node



Important

Follow the SUSE OpenStack Cloud documentation's [Deploying Nova](#) instructions.


```
# The name of the cluster.
cluster_name: 'midonet'

...

# Addresses of hosts that are deemed contact points.
seed_provider:
  - class_name: org.apache.cassandra.locator.SimpleSeedProvider
    parameters:
      - seeds: "nsdb1,nsdb2,nsdb3"
```

b. Node-specific Configuration

i. NSDB Node 1

Edit the `/etc/cassandra/conf/cassandra.yaml` file to contain the following:

```
# Address to bind to and tell other Cassandra nodes to connect to.
listen_address: nsdb1

...

# The address to bind the Thrift RPC service.
rpc_address: nsdb1
```

ii. NSDB Node 2

Edit the `/etc/cassandra/conf/cassandra.yaml` file to contain the following:

```
# Address to bind to and tell other Cassandra nodes to connect to.
listen_address: nsdb2

...

# The address to bind the Thrift RPC service.
rpc_address: nsdb2
```

iii. NSDB Node 3

Edit the `/etc/cassandra/conf/cassandra.yaml` file to contain the following:

```
# Address to bind to and tell other Cassandra nodes to connect to.
listen_address: nsdb3

...

# The address to bind the Thrift RPC service.
rpc_address: nsdb3
```

3. Edit the service's init script

On installation, the `/var/run/cassandra` directory is created, but because it is located on a temporary file system it will be lost after system reboot. As a result it is not possible to stop or restart the Cassandra service anymore.

To avoid this, edit the `/etc/init.d/cassandra` file to create the directory on service start:

```
case "$1" in
    start)
        # Cassandra startup
        echo -n "Starting Cassandra: "
        mkdir -p /var/run/cassandra
        chown cassandra:cassandra /var/run/cassandra
        su $CASSANDRA_OWNRR -c "$CASSANDRA_PROG -p $pid_file" > $log_file
2>&1
        retval=$?
[...]
```

4. Enable and start Cassandra

```
# service cassandra start
```

5. Verify Cassandra Operation

After installation of all nodes has been completed, verify that Cassandra is operating properly.



Important

If Cassandra fails to start and prints a "buffer overflow" error message in its log file, you may try associating `127.0.0.1` with the hostname in `etc/hosts` (so that `hostname -i` will show `127.0.0.1`). This may solve the Cassandra start problem.

A basic check can be done by executing the `nodetool status` command. This will reply with UN (Up / Normal) in the first column if the servers are running in a non-error state:

```
$ nodetool -host 127.0.0.1 status
[...]
```

Status=Up/Down

|/ State=Normal/Leaving/Joining/Moving

--	Address	Load	Tokens	Owns	Host ID
	Rack				
UN	192.0.2.1	123.45 KB	256	33.3%	
	11111111-2222-3333-4444-555555555555				rack1
UN	192.0.2.2	234.56 KB	256	33.3%	
	22222222-3333-4444-5555-666666666666				rack1
UN	192.0.2.3	345.67 KB	256	33.4%	
	33333333-4444-5555-6666-777777777777				rack1

Controller Node

MidoNet Cluster Installation

1. Install MidoNet Cluster package

```
# zypper install midonet-cluster
```

2. Set up mn-conf

Edit `/etc/midonet/midonet.conf` to point `mn-conf` to the ZooKeeper cluster:

```
[zookeeper]  
zookeeper_hosts = nsdb1:2181,nsdb2:2181,nsdb3:2181
```

3. Configure access to the NSDB

This step needs to happen only once, it will set up access to the NSDB for the MidoNet Cluster and Agent nodes.

Run the following command to set the cloud-wide values for the ZooKeeper and Cassandra server addresses:

```
$ cat << EOF | mn-conf set -t default
zookeeper {
    zookeeper_hosts = "nsdb1:2181,nsdb2:2181,nsdb3:2181"
}

cassandra {
    servers = "nsdb1,nsdb2,nsdb3"
}
EOF
```

Run the following command to set the Cassandra replication factor:

```
$ echo "cassandra.replication_factor : 3" | mn-conf set -t default
```

4. Configure Keystone access

This step needs to happen only once, it will set up access to Keystone for the MidoNet Cluster node(s).

This step will configure the local Midonet Cluster node to be able to use Keystone.

```
$ cat << EOF | mn-conf set -t default
cluster.auth {
    provider_class = "org.midonet.cluster.auth.keystone.KeystoneService"
    admin_role = "admin"
    keystone.tenant_name = "admin"
    keystone.admin_token = "ADMIN_TOKEN"
    keystone.host = controller
    keystone.port = 35357
}
EOF
```

5. Start the MidoNet Cluster

```
# systemctl enable midonet-cluster.service
# systemctl start midonet-cluster.service
```

MidoNet CLI Installation

1. Install MidoNet CLI package

```
# zypper install python-midonetclient
```

2. Configure MidoNet CLI

Create the ~/.midonetr.c file and edit it to contain the following:

```
[cli]
api_url = http://controller:8080/midonet-api
username = admin
password = ADMIN_PASS
project_id = admin
```

Midolman Installation

The *MidoNet Agent (Midolman)* has to be installed on all nodes where traffic enters or leaves the virtual topology, in this guide this are the **gateway1**, **gateway2** and **compute1** nodes.

1. Install Midolman package

```
# zypper install midolman
```

2. Set up mn-conf

Edit `/etc/midolman/midolman.conf` to point mn-conf to the ZooKeeper cluster:

```
[zookeeper]
zookeeper_hosts = nsdb1:2181,nsdb2:2181,nsdb3:2181
```

3. Configure access to the NSDB for all agents

This step needs to happen only once, it will set up access to the NSDB for all MidoNet Agent nodes.

Run the following command to set the cloud-wide values for the ZooKeeper and Cassandra server addresses:

```
$ cat << EOF | mn-conf set -t default
zookeeper {
    zookeeper_hosts = "nsdb1:2181,nsdb2:2181,nsdb3:2181"
}

cassandra {
    servers = "nsdb1,nsdb2,nsdb3"
}
EOF
```

Run the following command to set the Cassandra replication factor:

```
$ echo "cassandra.replication_factor : 3" | mn-conf set -t default
```

4. Configure resource usage

Run these steps on **each agent host** in order to configure resource usage.



Important

For production environments the **large** templates are strongly recommended.

a. Midolman resource template

Run the following command to configure the Midolman resource template:

```
$ mn-conf template-set -h local -t TEMPLATE_NAME
```

Replace **TEMPLATE_NAME** with one of the following templates:

```
agent-compute-large
agent-compute-medium
agent-gateway-large
agent-gateway-medium
default
```

b. Java Virtual Machine (JVM) resource template

Replace the default `/etc/midolman/midolman-env.sh` file with one of the below to configure the JVM resource template:

```
/etc/midolman/midolman-env.sh.compute.large  
/etc/midolman/midolman-env.sh.compute.medium  
/etc/midolman/midolman-env.sh.gateway.large  
/etc/midolman/midolman-env.sh.gateway.medium
```

5. Start Midolman

```
# systemctl enable midolman.service  
# systemctl start midolman.service
```

MidoNet Host Registration

1. Launch MidoNet CLI

```
$ midonet-cli  
midonet>
```

2. Create tunnel zone

MidoNet supports the Virtual Extensible LAN (VXLAN) and Generic Routing Encapsulation (GRE) protocols to communicate to other hosts within a tunnel zone.

To use the VXLAN protocol, create the tunnel zone with type 'vxlan':

```
midonet> tunnel-zone create name tz type vxlan  
tzone0
```

To use the GRE protocol, create the tunnel zone with type 'gre':

```
midonet> tunnel-zone create name tz type gre  
tzone0
```



Important

Make sure to allow GRE/VXLAN traffic for all hosts that belong to the tunnel zone. For VXLAN MidoNet uses UDP port 6677 as default.

1. Add hosts to tunnel zone

```
midonet> list tunnel-zone  
tzone tzone0 name tz type vxlan  
  
midonet> list host  
host host0 name controller alive true  
host host1 name gateway1 alive true  
host host2 name gateway2 alive true  
host host3 name compute1 alive true  
  
midonet> tunnel-zone tzone0 add member host host0  
address ip_address_of_host0  
zone tzone0 host host0 address ip_address_of_host0  
  
midonet> tunnel-zone tzone0 add member host host1  
address ip_address_of_host1  
zone tzone0 host host1 address ip_address_of_host1
```

```
midonet> tunnel-zone tzone0 add member host host2
address ip_address_of_host2
zone tzone0 host host2 address ip_address_of_host2

midonet> tunnel-zone tzone0 add member host host3
address ip_address_of_host3
zone tzone0 host host3 address ip_address_of_host3
```


5. Initial Network Configuration

1. To create the external network

Use the following command to create the external network:

```
$ neutron net-create ext-net --router:external
```

6. Edge Router Setup

Prior to v5.0, with Neutron, you could set up the gateway in only one way, which was to have a special singleton gateway router called the Provider Router created implicitly when an external network was created in Neutron. The provider router sits at the edge of the cloud and interfaces with the uplink router. The Provider Router is where BGP was typically configured. The biggest limitation of this approach was that it took away the scenario in which you wanted to have an L2 network at the edge instead of a router. Another limitation was that only one such router could exist for the entire cloud.

These limitations are removed in v5.0, where you could design your gateway to be either L2 network or router with as many routers as you wish, all using the Neutron API.

There are two main changes:

Edge Router

The Provider Router is no longer implicitly created upon the external network creation. Instead, the edge gateway routers, called the Edge Routers, are created explicitly using standard Neutron API. With this approach, multiple Edge Routers can be created, and they are optional.

Gateway Virtual Topology

In the previous model, the Provider Router was connected directly to the tenant routers, with the external networks hanging off of the Provider Router.

In the new model, the external networks exist between the edge and the tenant routers.

To create the gateway topology issue the following Neutron commands.

Create a standard neutron router:

```
neutron router-create <EDGE_ROUTER_NAME>
```

Attach the edge router to an external network:

```
neutron router-interface-add <EDGE_ROUTER_ID> <EXT_SUBNET_ID>
```

Create a special network called uplink network, representing the physical network outside of the cloud:

```
neutron net-create <UPLINK_NET_NAME> --tenant_id admin --  
provider:network_type uplink
```

Create a subnet for the uplink network matching the CIDR used in the uplink network (could just be /30 if linked directly to another router):

```
neutron subnet-create --tenant_id admin --disable-dhcp --name  
<UPLINK_SUBNET_NAME> <UPLINK_NET_NAME> <CIDR>
```

Create a port on the uplink network with a specific IP that you want to use and the binding details so that this virtual port gets bound to a specific NIC on the gateway host:

```
neutron port-create <UPLINK_NET_ID> --binding:host_id <HOST_NAME> --  
binding:profile type=dict interface_name=<INTERFACE_NAME> --fixed-ip  
ip_address=<IP_ADDR>
```

Attach the uplink port to the Edge Router:

```
neutron router-interface-add <EDGE_ROUTER_ID> port=<UPLINK_PORT_ID>
```


8. Further Steps

MidoNet installation and integration into OpenStack is completed.



Note

Consult the **Operations Guide** for further instructions on operating MidoNet.